

Concurrency and Multithreading in Java


Moeen Aali

Sharif University of Technology

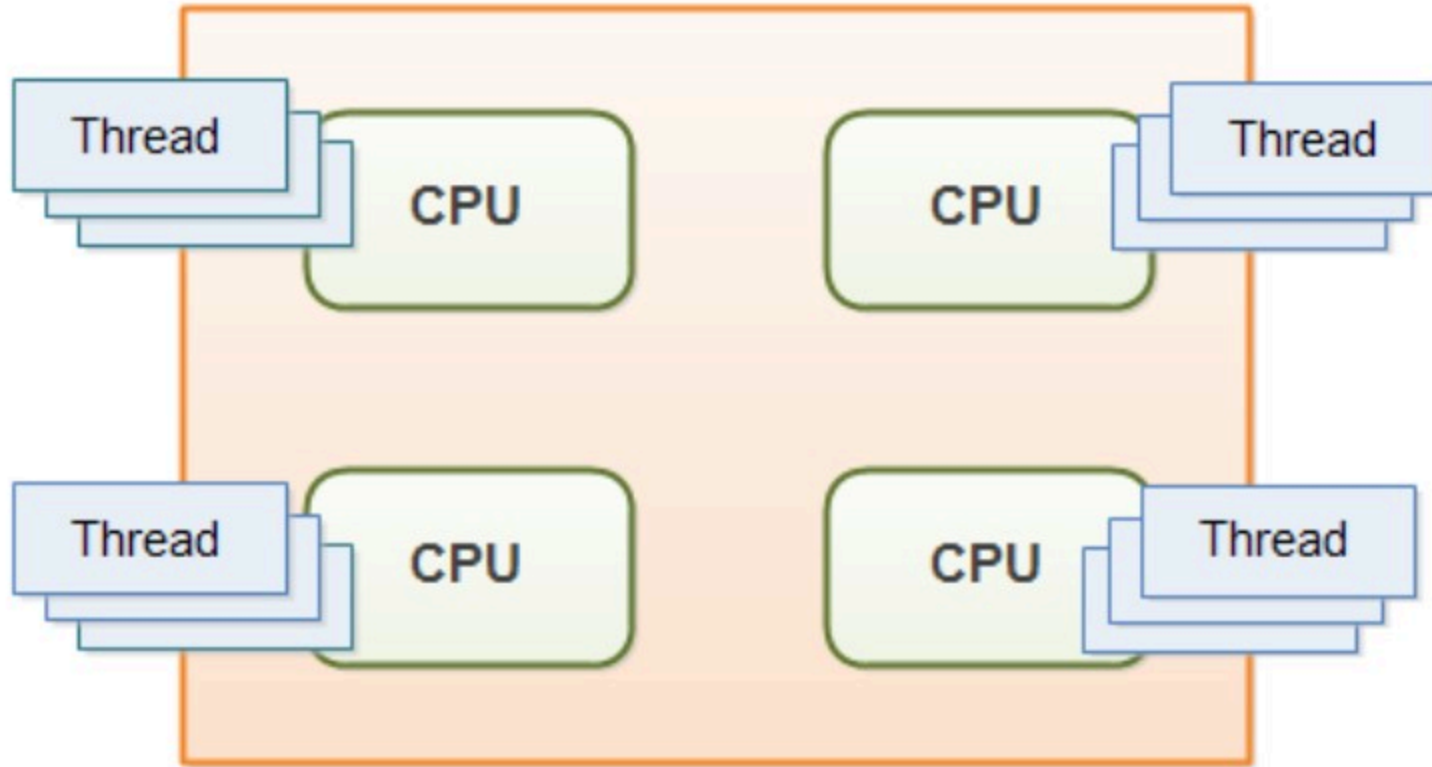
CE244 - Advanced Programming - Fall 2024

Mr. Hamidreza Hosseinkhani

Concurrency vs. Parallelism



Concurrency vs. Parallelism



- Even without the ability to execute in parallel, the ability to concurrency is useful.


Thread

When a Java program runs:

- A thread is created to execute the `main()` method.

Key Points:

- The program can create new threads and execute them.
- Multiple threads run **concurrently**.
 - This can sometimes mean running **in parallel**.



Thread: Create

Two primary ways to define the behavior of a new thread:

1. Create a class that **extends** `java.lang.Thread`.
2. Create a class that **implements** the `java.lang.Runnable` interface.
 - Implement the `run` method in the new class.

Thread: Create a class that extends `java.lang.Thread`

```
class MyThread extends Thread {  
    @Override  
    public void run() {  
        System.out.println("Hello");  
        System.out.println("Bye");  
    }  
}
```

```
public class Main{  
    public static void main(String[] args) {  
        System.out.println("Salam");  
        MyThread t = new MyThread();  
        t.start();  
        System.out.println("Khodahafez");  
    }  
}
```

Thread: Implements the `java.lang.Runnable` interface

```
class MyRunnable implements Runnable{  
    @Override  
    public void run() {  
        System.out.println("Hello");  
        System.out.println("Bye");  
    }  
}
```

```
Thread t = new Thread(new MyRunnable());  
t.start();
```

Which Approach is Better?

Extending `Thread`:

- Simpler implementation.
- However, limits the class from **inheriting** from other classes.
 - A class cannot extend multiple classes in Java.

Implementing `Runnable`:

- Provides more **flexibility**.
- Allows the class to extend a different class if needed.
- Commonly used in practice.

Run vs. Start

Why Implement `run` but Call `start`?

- The `start` method is a **special method** in the `Thread` class.
 - It performs low-level system operations to create a new thread.
 - Calls the `run` method **inside the new thread**.
- Directly calling `run`:
 - Executes the method like a **normal function call**.
 - Does **not** create a new thread.

Thread Methods

- For every thread that executes, an object of the `Thread` class is created.
- Methods of the `Thread` object provide functionalities for the corresponding thread.

Some Methods of the `Thread` Class

- Common instance methods:
 - `run`, `start`, `getId`, `setPriority`, `setDaemon`, ...

Thread Methods

Some Static Methods

- `currentThread` : Returns the currently executing thread.
- `sleep` : Puts the currently executing thread to sleep for a specified time.
 - This pauses its execution for the given duration and then resumes it.
- `join` : Sometimes, it is necessary for one thread to complete its work before another section of code can proceed.

```
Thread myThread = new MyThread();  
myThread.start();  
func1();  
myThread.join();  
func2();
```

Thread Priority

- The priority of a thread is adjustable.
- It can be changed using the `setPriority` method.

Priority Levels

- Thread priority is a number between `1` and `10` that indicates the thread's importance.
- Higher-priority threads are given more execution time by the operating system.
 - Threads with higher priority are allocated more CPU time.

```
MyThread thread = new MyThread();  
thread.setPriority(Thread.MAX_PRIORITY);  
thread.start();
```

```
Thread.MIN_PRIORITY = 1;  
Thread.NORM_PRIORITY = 5;  
Thread.MAX_PRIORITY = 10;
```

Daemon Threads

- A special type of thread that runs in the background.
- Typically, they provide services to other threads and do not have independent meaning or purpose.
- The Garbage Collector (`GC`) is an example of a daemon thread.
- If only daemon threads remain alive in a program and all regular threads have finished, the JVM terminates the daemon threads and the program ends.
- Use the `setDaemon(boolean on)` method to mark a thread as either daemon (`true`) or regular (`false`).

```
MyThread thread = new MyThread();  
thread.setDaemon(true);  
thread.start();
```

Quiz: How many threads does this program have?

```
class T extends Thread {
    public void run() {
        for (int i = 1; i <= 100; i++)
            System.out.println(i);
    }
}
class R implements Runnable{
    public void run() {
        for (char c = 'A'; c < 'Z'; c++)
            System.out.println(c);
    }
}
public class Threading{
    public static void main(String[] args) {
        new Thread(new R()).start();
        new T().start();
        new Thread(new R()).start();
        new T().start();
    }
}
```

Program's memory

Recap

- In a program's memory, there are distinct sections like:
 - **Stack**: Stores local variables.
 - **Heap**: Stores objects.

Thread-Specific Memory

- Each thread has its **own stack**.
 - For example, if two threads call the same method, each thread will have a separate memory space for the method's local variables in their respective stacks.

Shared Memory

- All threads share the **Heap** memory.
 - Multiple threads can access and use the same object in the Heap.

Critical Sections in Multithreading

- Threads can simultaneously access and use shared objects.
- This can lead to problems, such as:
 - One thread modifying an object while another thread is also changing it.
 - One thread working on a file while another thread closes it.

Critical Sections

- A **critical section** is a part of the program that should not be executed by multiple threads at the same time.
- If one thread enters a critical section:
 - No other thread should be allowed to enter the same section until the first thread finishes its execution.

Key Behavior

- The execution of a second thread should pause until the first thread exits the critical section.

Critical Sections

- Java automatically handles the pausing of threads when necessary.
 - When one thread is entering a critical section, other threads attempting to enter the same section will be blocked if it is already being executed.

Locking Critical Sections

- Java provides mechanisms to define and manage critical sections.
- The programmer must specify which parts of the program are critical and define the conditions for entering them.

Thread Locking

- When a thread enters a critical section, it acquires a **lock**.
 - If another thread has already acquired the lock, the thread must wait until the lock is released.
- When exiting the critical section, the thread releases the lock it has acquired.
- The programmer determines which lock is required for each critical section.

Synchronized

- In Java, any object can be used as a lock for entering a critical section.

synchronized Method

- When a **synchronized** method is invoked on an object, the thread tries to acquire the lock for that object before entering the method.
 - This means it attempts to acquire the **this** lock for the object.
 - There is only one lock for each object.

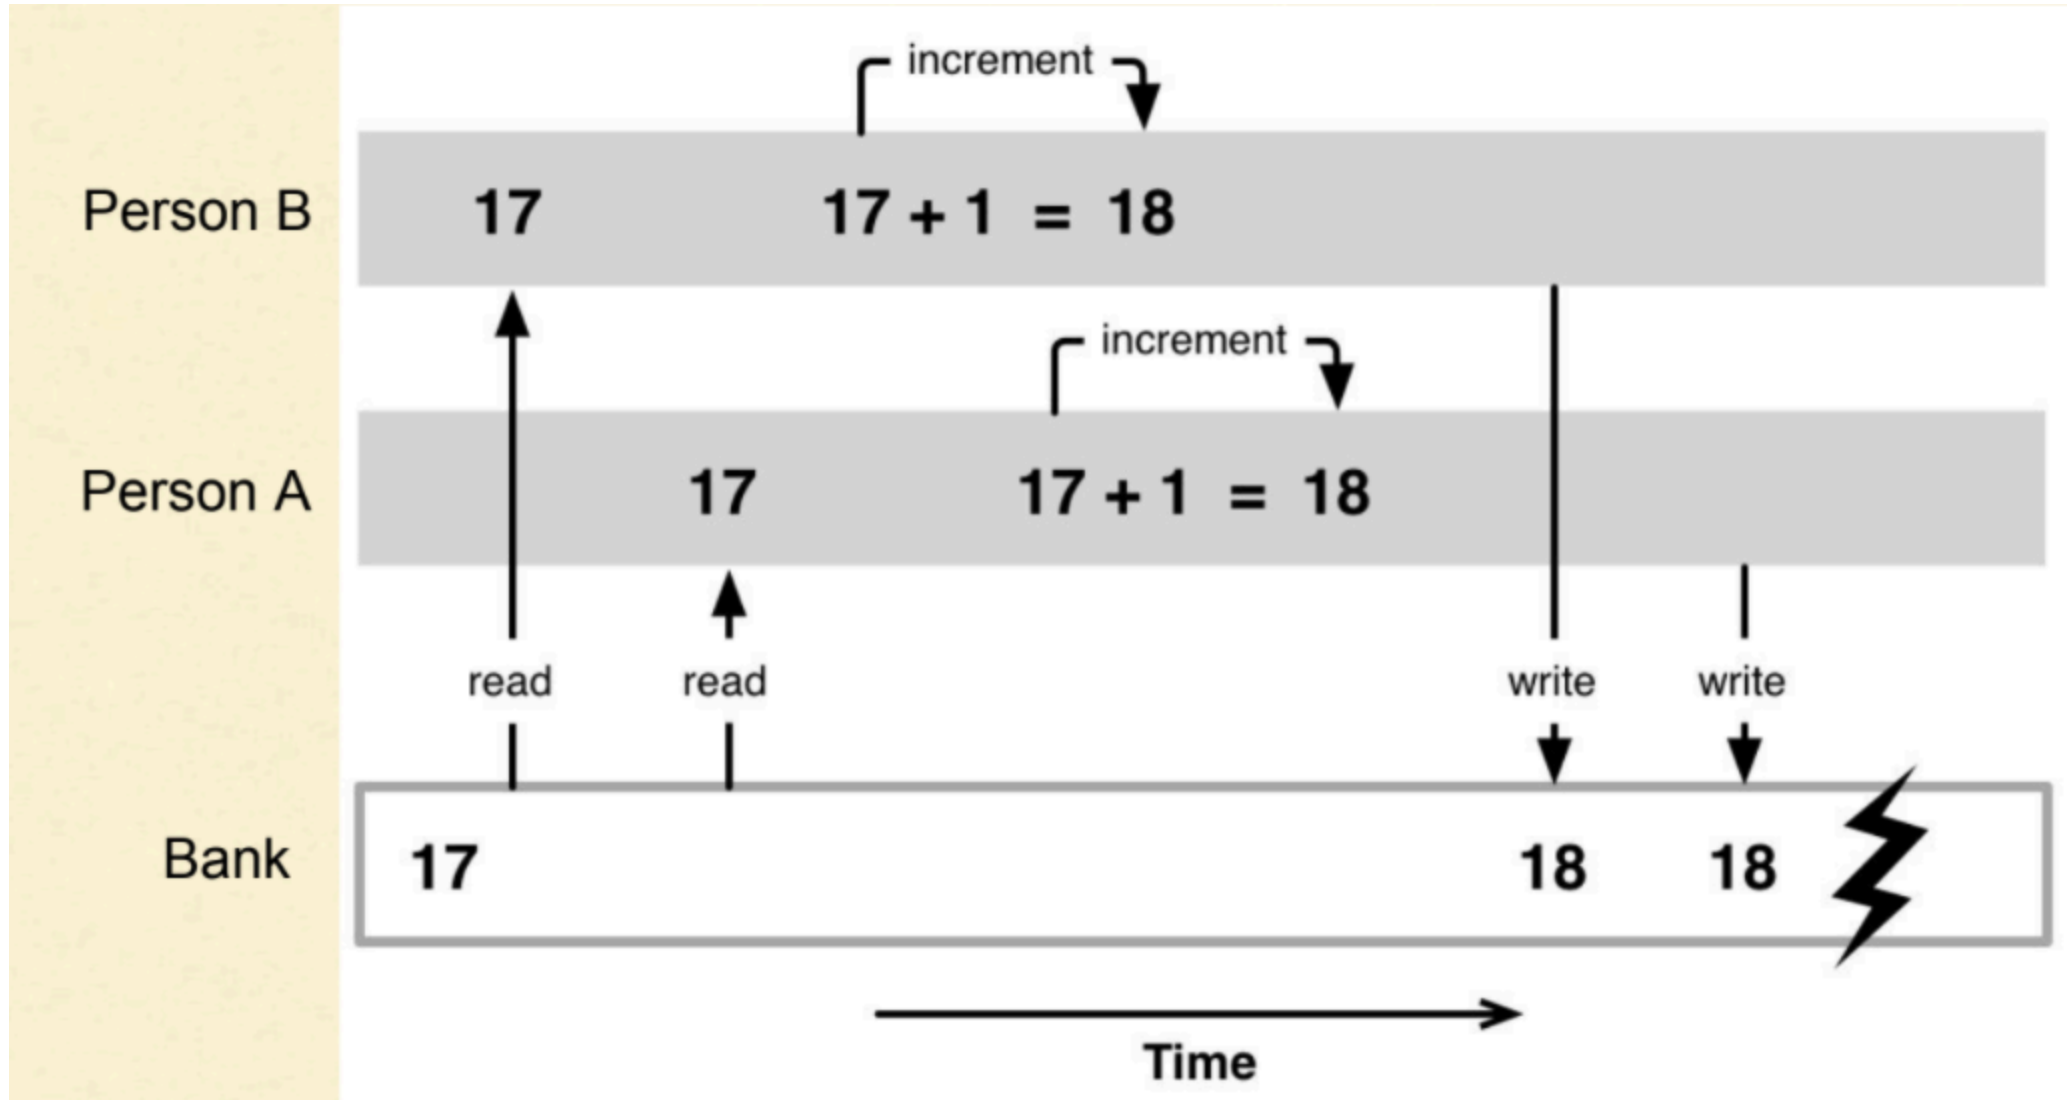
Locking Behavior

- When a **synchronized** method is executing:
 - No other **synchronized** method on the same object can start execution simultaneously.
 - This is because, until the method finishes, no other thread can acquire the **this** lock.

Synchronized

```
public class BankAccount {  
    private float balance;  
    public synchronized void deposit(float amount) {  
        balance += amount;  
    }  
    public synchronized void withdraw(float amount) {  
        balance -= amount;  
    }  
}
```

Synchronized: Race Condition



Synchronized

Using `synchronized` Block with an Object

- You can create a `synchronized` block and specify an object to lock.

Example:

- Two different threads can enter the synchronized block simultaneously, **only if** the objects they are synchronizing on are different.
 - This allows fine-grained control over which threads can access specific blocks of code concurrently, based on the object used for synchronization.

Synchronized

These two definitions are approximately equivalent for the g method:

```
void g() {  
    synchronized(this){  
        h();  
    }  
}
```

```
synchronized void g() {  
    h();  
}
```

- A non-static synchronized method may be executed by two threads at the same time, provided they are called on two different objects.

Inter-thread Communication

`wait` and `notify` Methods in Java

- Sometimes, it is necessary for two threads to interact with each other.
 - A thread may need to wait until another thread notifies it.

Behavior:

- When a thread calls the `wait` method on an object:
 - It pauses until another thread calls the `notify` method on the same object, which wakes up the waiting thread.
- You can use newer methods like `Synchronizer`.

Inter-thread Communication

- The `wait` and `notify` methods can only be called on an object `x` if they are inside a `synchronized(X)` block.
 - A thread must acquire the lock for the object `x` before calling `wait` or `notify` on it.

Key Notes:

- If a thread tries to call `wait` or `notify` on an object without holding its lock, an `IllegalMonitorStateException` will be thrown.
- When a thread calls `x.wait()`, the lock on `x` is released immediately.
 - This allows other threads to enter the `synchronized(X)` block and call `x.notify()` to wake up the waiting thread.

```
synchronized (obj) {  
    obj.notify();  
}  
  
synchronized void f() {  
    wait();  
}
```


Inter-thread Communication

- Multiple threads can be waiting on the same object.
 - Each object maintains a list of threads that are waiting for it.

Waking Up Threads

- When the `notify` method is called on an object, **one** of the waiting threads will be awakened and continue its execution.
- The `notifyAll` method will wake up **all** waiting threads on the object.

Timeout with `wait`

- The `wait` method can specify a maximum waiting time.
 - For example, `wait(100)` means the thread will wait for up to **100 milliseconds** before waking up, even if the `notify` method has not been called by another thread during this time.

Thread Interrupt

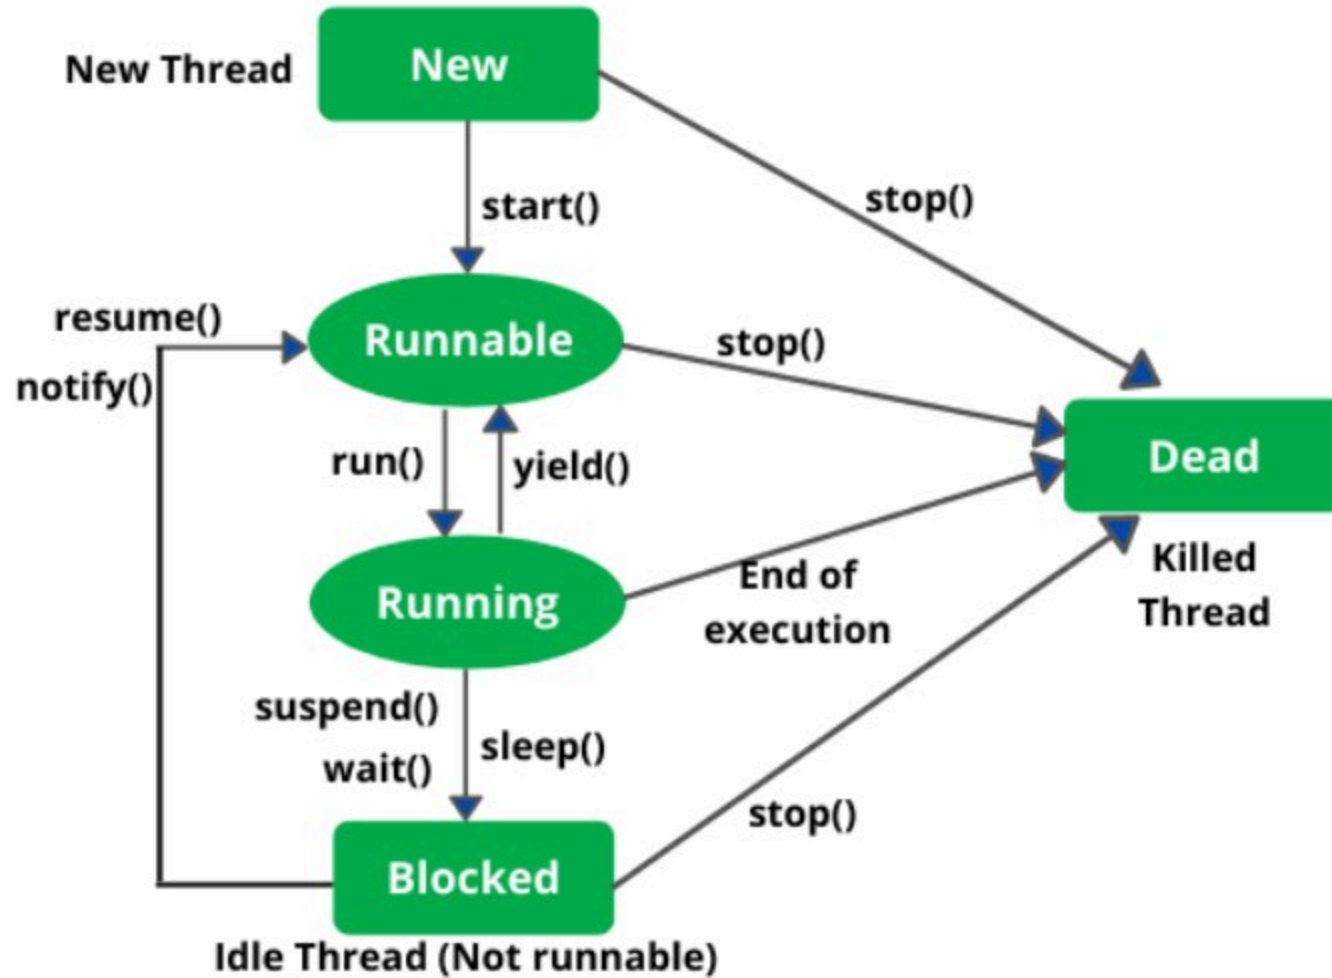
- Interrupts are used to **stop or pause** the execution of a thread.
- A thread can handle an interrupt by checking if it has been interrupted using methods like:
 - `Thread.interrupted()`
 - `isInterrupted()`

Why Use Interrupts?

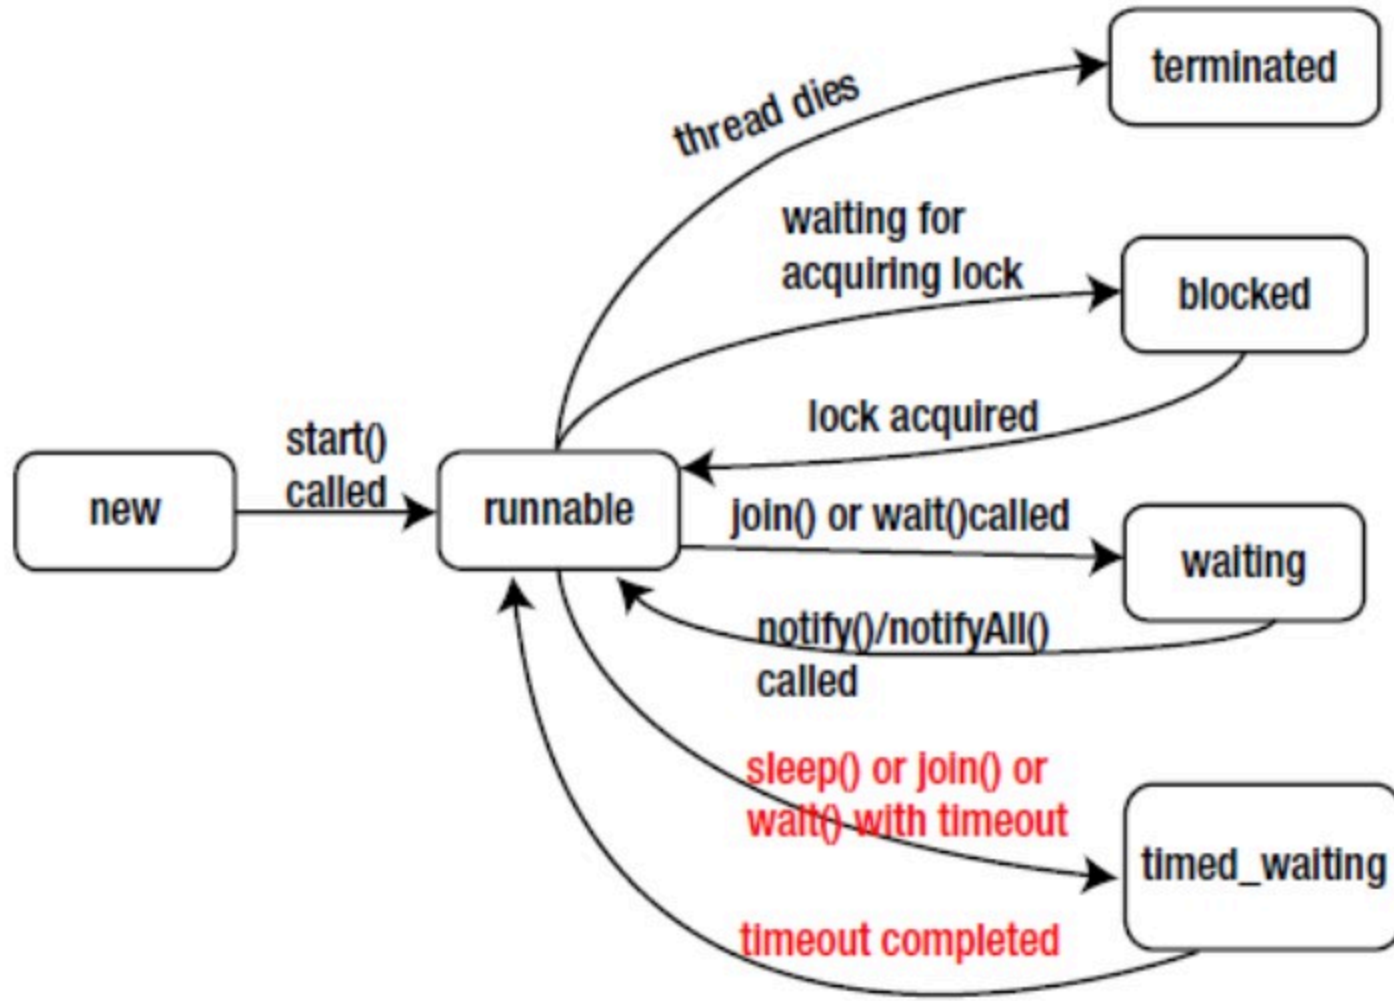
- Graceful termination of threads.
- Managing thread lifecycle effectively, especially in multi-threaded applications.

```
Thread thread = new Thread(() -> {  
    while (!Thread.currentThread().isInterrupted()) {  
        // Do work  
    }  
});  
  
thread.interrupt();
```

Thread Life Cycle



Thread States



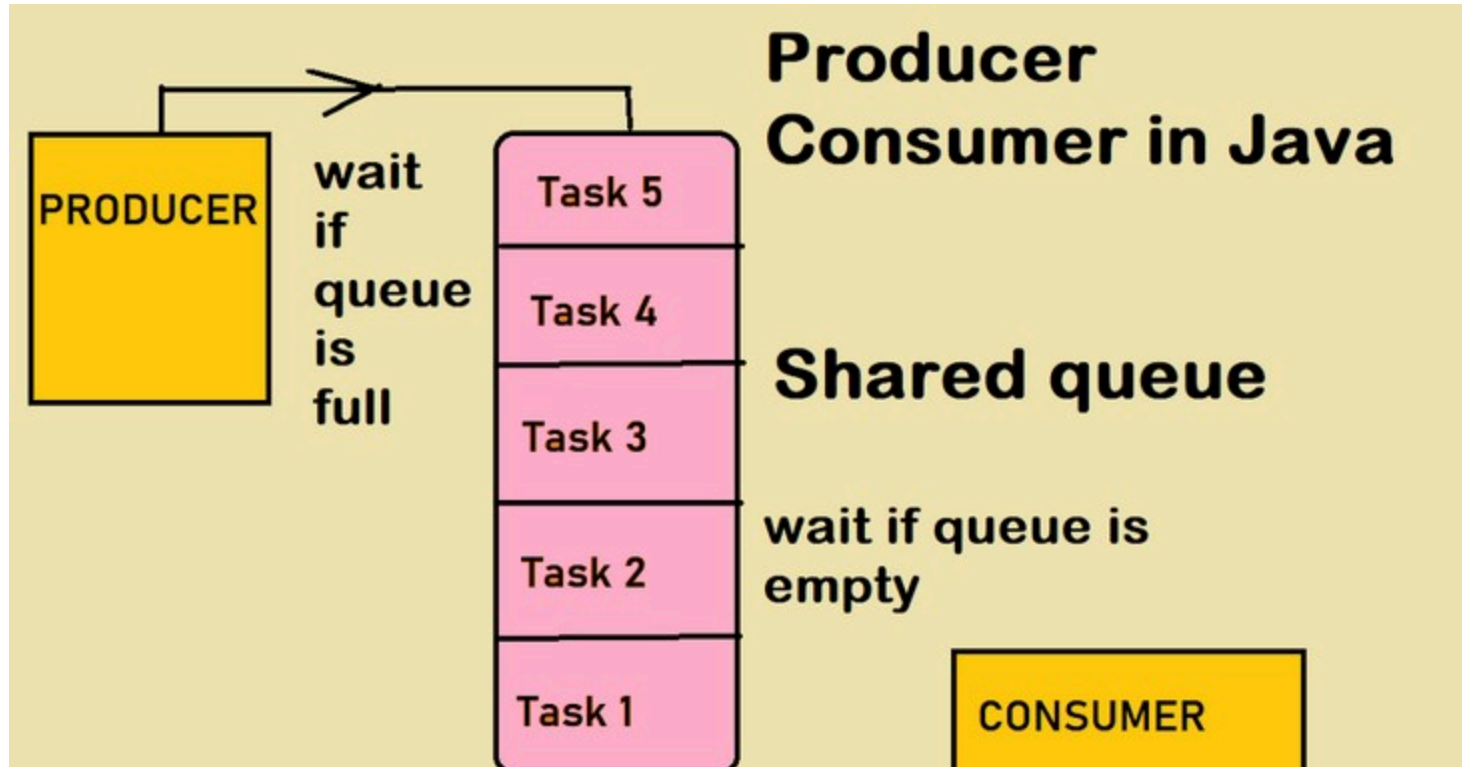
Thread States

- The `getState()` method for each Thread object returns the state of that thread.

```
public enum State {  
    NEW,  
    RUNNABLE,  
    BLOCKED,  
    WAITING,  
    TIMED_WAITING,  
    TERMINATED;  
}
```

Example: Producer/Consumer

- a classic example of a multi-process synchronization problem



Thread-Safe Classes

- A **thread-safe** class ensures that objects of that class can be used concurrently by multiple threads without causing issues.
- No need for additional synchronization (`synchronized`) or locks when using these objects across multiple threads.

Class	Description
<code>Vector</code>	Thread-safe version of <code>ArrayList</code>
<code>ConcurrentHashMap</code>	Thread-safe version of <code>HashMap</code>
<code>StringBuffer</code>	Thread-safe version of <code>StringBuilder</code>
<code>ArrayBlockingQueue</code>	Thread-safe version of <code>Queue</code>

Thread-Safe Classes

- No explicit synchronization required for safe use in multi-threaded environments.
- Thread safety mechanisms are implemented internally in these classes.
- Can be used across multiple threads without worrying about data integrity.
- Always Immutable objects are Thread-safe.
- We shouldn't use Vector instead of ArrayList at any time.

Thread-Safe Queue

- Key methods:
 - `put` : Adds an element to the queue.
 - `take` : Removes and retrieves an element from the queue.

Blocking Behavior

- Threads are **blocked** as needed during reading or writing:
 - **Reading**: If the queue is empty, the thread waits until an element is available.
 - **Writing**: If the queue is full, the thread waits until space becomes available.

Implementation	Description
<code>ArrayBlockingQueue</code>	Fixed-size queue implemented using an array.
<code>LinkedBlockingQueue</code>	Dynamically-sized queue implemented using a linked list.

Semaphore

- A **Semaphore** is a thread synchronization construct.
- Used to control access to a shared resource with a fixed number of permits.
- **Key Methods:**
 - `acquire()` : Acquires a permit, blocking if none are available.
 - `release()` : Releases a permit, increasing the number of available permits.
- Can be used to:
 - Limit concurrent access to a resource.
 - Implement simple thread signaling.

Semaphore

```
Semaphore semaphore = new Semaphore(2);
Runnable task = () -> {
    try {
        System.out.println(Thread.currentThread().getName() + " trying to acquire...");
        semaphore.acquire(); // Acquire a permit
        System.out.println(Thread.currentThread().getName() + " acquired a permit.");
        Thread.sleep(2000); // Simulate work
    } catch (InterruptedException e) {
        Thread.currentThread().interrupt();
    } finally {
        System.out.println(Thread.currentThread().getName() + " releasing permit.");
        semaphore.release(); // Release the permit
    }
};
for (int i = 0; i < 5; i++) {
    new Thread(task).start();
}
```

Atomic Classes

- Atomic classes are part of `java.util.concurrent.atomic` package.
- Provide low-level thread-safe operations on single variables.
- Avoid the need for explicit synchronization.

Key Features of Atomic Classes

1. **Atomicity:** Operations are thread-safe without explicit locks.
2. **Performance:** Lightweight compared to synchronization blocks.
3. **Convenience Methods:**
 - `get()` and `set()`
 - `incrementAndGet()` / `decrementAndGet()`
 - `compareAndSet(expectedValue, newValue)`

Atomic Classes

Common Atomic Classes

- `AtomicInteger` : Atomic operations on `int` .
- `AtomicLong` : Atomic operations on `long` .
- `AtomicBoolean` : Atomic operations on `boolean` .
- `AtomicReference<V>` : Atomic operations on references.

Why Use Atomic Classes?

- Simplifies code for multithreaded applications.
- Reduces the chance of errors due to improper synchronization.
- Ideal for counters, flags, and state management.

Atomic Classes: AtomicInteger

Why Use Atomic Classes?

```
import java.util.concurrent.atomic.AtomicInteger;

public class AtomicExample {
    public static void main(String[] args) {
        AtomicInteger counter = new AtomicInteger(0);
        int newValue = counter.incrementAndGet();
        System.out.println("Counter: " + newValue);
        boolean updated = counter.compareAndSet(1, 10);
        System.out.println("Updated: " + updated);
        System.out.println("Counter: " + counter.get());
    }
}
```

ReentrantReadWriteLock

```
List<Double> list= new LinkedList<>();
ReadWriteLock lock = new ReentrantReadWriteLock();

class Reader extends Thread{
    public void run() {
        lock.readLock().lock();
        System.out.println(list.get(0));
        lock.readLock().unlock();
    }
}

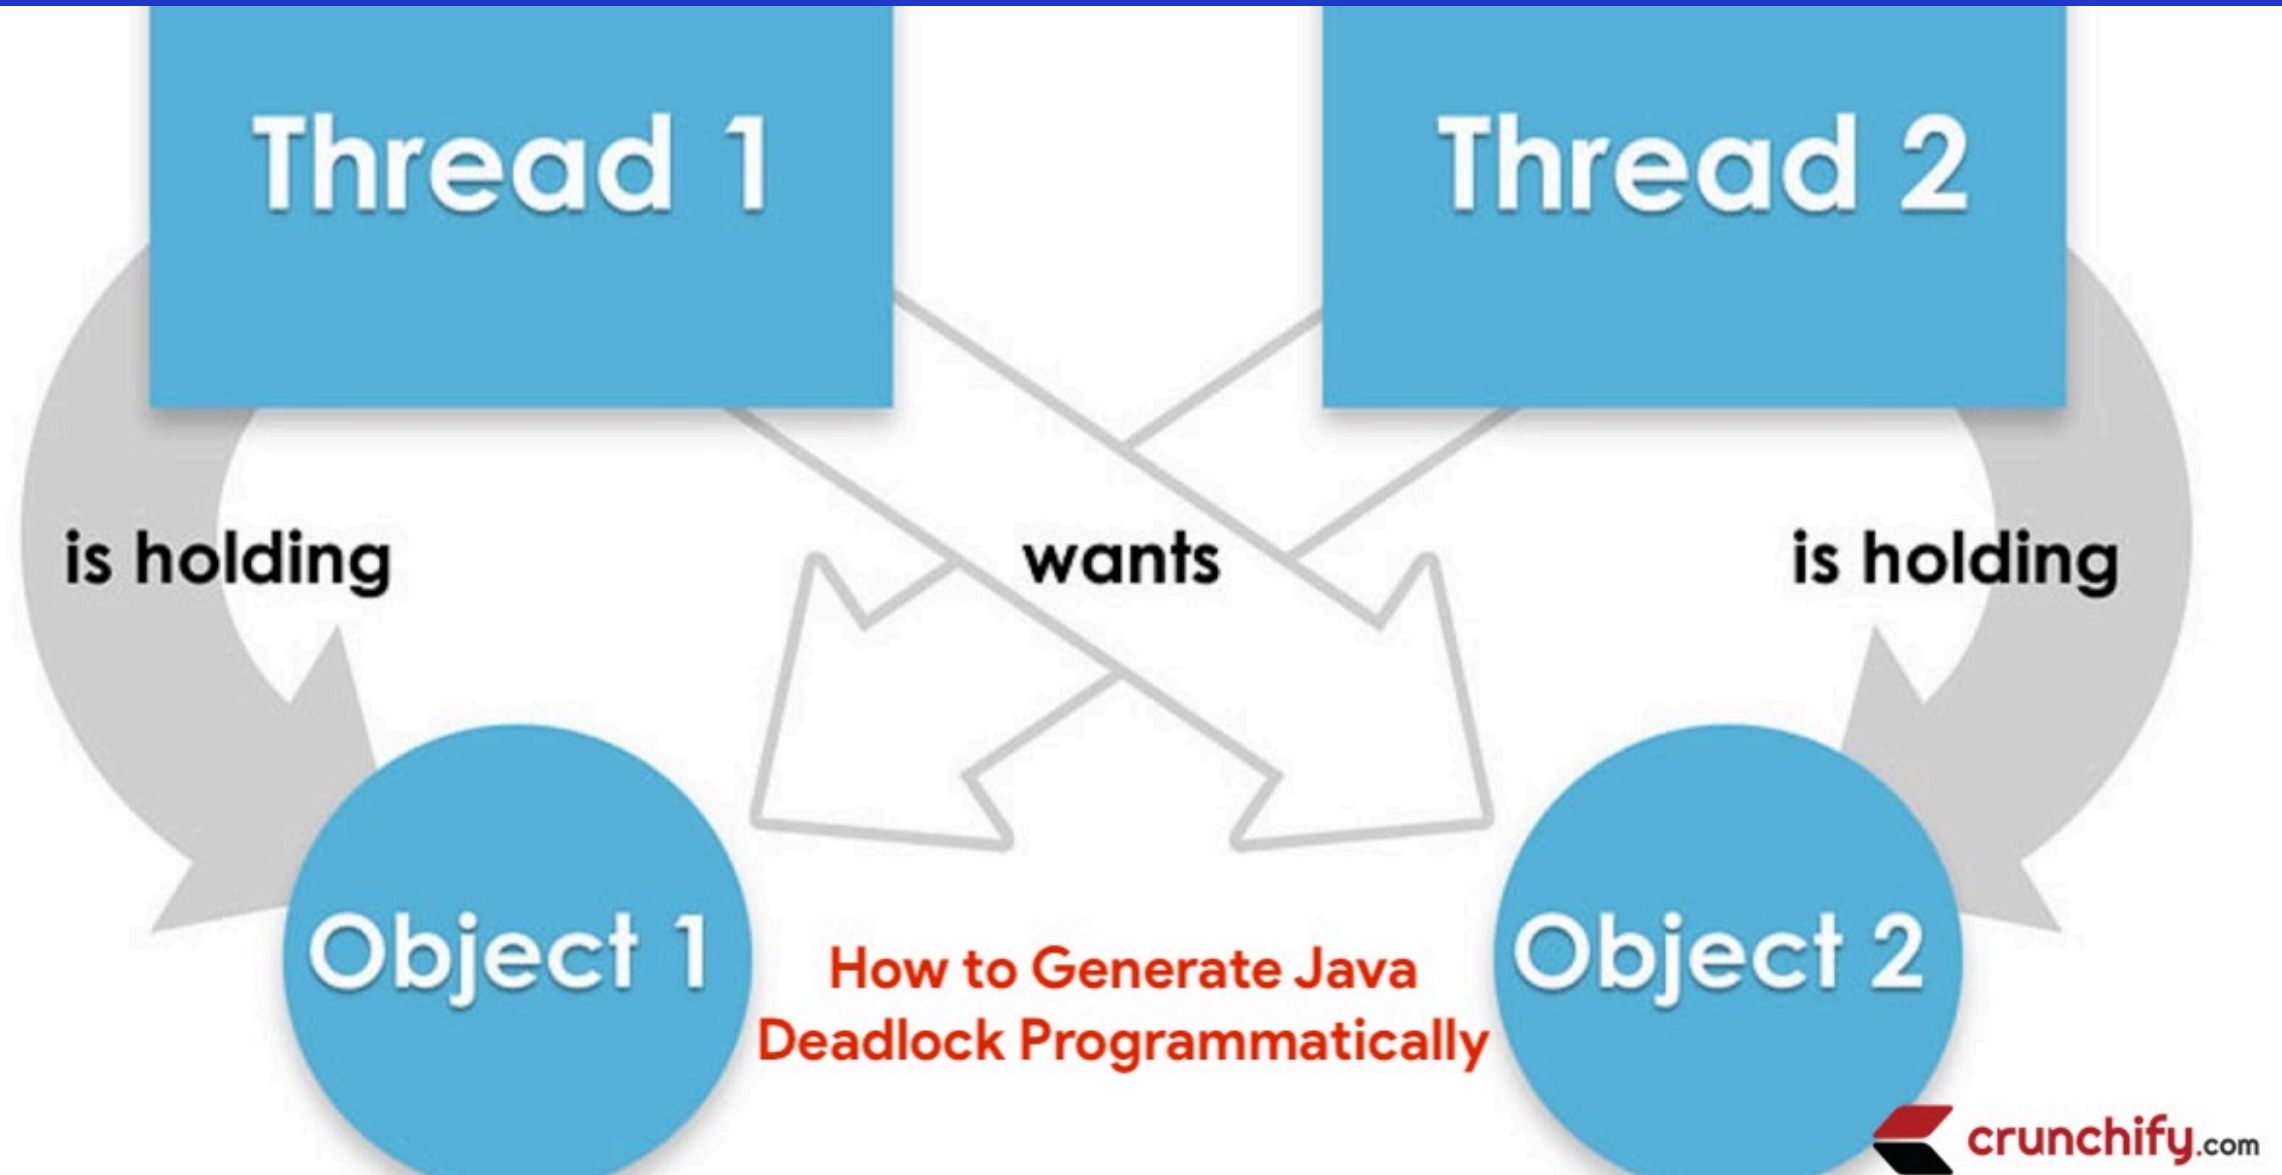
class Writer extends Thread{
    public void run() {
        lock.writeLock().lock();
        list.add(0, Math.random());
        lock.writeLock().unlock();
    }
}
```

Deadlock

- Occurs when two or more threads are waiting indefinitely for resources held by each other.
- **Conditions for Deadlock:**
 - i. Mutual Exclusion
 - ii. Hold and Wait
 - iii. No Preemption
 - iv. Circular Wait

```
synchronized (resource1) {  
    synchronized (resource2) {  
        // Deadlock can occur if threads lock resources in opposite order  
    }  
}
```


Deadlock



References

- Java How to Program - Deitel & Deitel
- Java Cup